

Blueprint for a Particle Theory AI Community for AI-Enabled Discoveries

George T. Fleming¹, Daniel C. Hackett¹, James Halverson², T. J. Hobbs³, Stefan Hoeche¹, Brandon Kriesten³, Konstantin T. Matchev⁴, Tony Menzo^{1,4}, Stephen Mrenna¹, Benjamin Nachman⁵, Prasanth Shyamsundar¹, David Shih⁶, Jesse Thaler⁷, Jure Zupan⁸

¹ Fermi National Accelerator Laboratory, ² Northeastern University, ³ Argonne National Laboratory,

⁴ University of Alabama, ⁵ SLAC National Laboratory, ⁶ Rutgers University,

⁷ Massachusetts Institute of Technology, ⁸ University of Cincinnati

September 3, 2026

ABSTRACT

Rapid advances in the capabilities and reach of artificial intelligence (AI) are changing research in particle physics. In this new era, AI is an innovative tool for theoretical physicists, while the interaction of AI and particle theory is itself a developing field. We argue that particle theory requires a community-scale effort to realize its full potential. We show how particle theory occupies a unique position in the AI for physics landscape, and describe the complementarity of AI and theory developments. We conclude with a series of recommendations on how to build an AI theory community for the Genesis Mission.

1 RESPONSE TO THE NATIONAL CHALLENGES

This white paper is a response to the Genesis Mission (GM)¹ from a group of particle theorists at universities and national laboratories. The GM aims to stimulate advances in Artificial Intelligence (AI), which includes Machine Learning (ML) and Large Language Models (LLMs). Two of the Science and Technology challenges identified by the Genesis Mission (GM) [1], Unifying Physics from Quarks to Cosmos and Discovering Quantum Algorithms with AI, are aligned with the particle theory community:

- to integrate the vast array of microscopic to cosmological data with theoretical knowledge to accelerate discovery.
- to accelerate the design and development of quantum circuits and algorithms that represent realistic quantum field theories (QFTs) and have provable quantum advantages.

Our purpose is to highlight the role particle theory plays in the GM and provide a vision how theorists can help realize the GM goals. This contribution is the theory counterpart to the experiment-oriented white paper [2], which sets out an AI-native research ecosystem for experimental particle physics. We take theory and experiment to be two wings of a single national HEP AI effort rather than parallel enterprises. However, we argue that the theory wing is distinct and valuable and requires its own community-wide representation.

¹See <https://www.energy.gov/documents/genesis-mission-science-and-technology-challenges>.

1.1 HOW DO AND WILL THEORETICAL PHYSICISTS USE AI

The GM aims to utilize the latest and quickly developing advances in AI to accelerate the pace of scientific discovery. Developments within just the last year include: large multimodal models (LMMs) that are pretrained to receive information in various modalities, e.g. text, images, video, audio; LLMs trained to serve as AI agents via tool calling; and, within just the last *few months*, “coding agents” that carry out nontrivial computational (science) tasks with minimal expert human guidance [3]. Theoretical physicists are positioned advantageously to use these latest developments in two broad ways:

Theory Building, which means building physical models based on some theoretical principles, usually encoded in the language of mathematics, and then performing calculations to test if they reproduce known physical results. Agentic AI can assist at every stage in this exploratory process at scales and speeds not previously possible [4–9].

Theory Discovery, a form of inverse problem, is analyzing large physical datasets to infer what underlying theories or models could have produced them and potentially discover new physical principles. Agentic AI has demonstrated its utility in exploring large datasets [10–14]. Furthermore, bespoke physics-specific LMMs can be pre-trained on large experimental or observational physics datasets and then mined for features that point to the most important aspects of the data that any good theoretical model must reproduce.

Particle theory serves an important and complementary role to experimental particle physics.² Theory building and discovery are its bread-and-butter. More specifically, particle theory has many specialized requirements that are not well suited to “off-the-shelf” AI applications, thereby driving the discovery of new ways to use AI. Theoretical physicists have contributed to many of the recent advances in fundamental AI, acquiring non-traditional skills in data science, modern applied statistics, AI/ML concepts and architectures. Moreover, particle theory can generate its own synthetic data that can be improved systematically and regenerated under different theoretical assumptions. Theory can both supply the data from which patterns are inferred and subject those inferences to systematic verification and validation, establishing representations and standards before they are needed at the HL-LHC, DUNE, the EIC, and beyond.

1.2 THE NEED FOR A THEORY AI COMMUNITY

Experimental collaborations have O(1000) members for a very good reason: a detector for a frontier experiment cannot be built and operated by a small group, and the detector provides valuable data, a joint asset shared for semi-independent research. To achieve their ends, the collaborations adopt structure and hierarchy that provide direction and oversight, enforce standards, evaluate the rigor and validity of scientific results and manage resources.

This scale of organization does not exist in the theory community. Even a well-funded research team at an academic institution working in theory might consist of at most a faculty advisor, a postdoctoral researcher and several graduate students (experimental groups can be only slightly larger, but they self-organize into larger research groups within collaborations). For the most part (see the USQCD example below), any larger scale organization of particle theory has not been viewed as necessary, and the independence of theoretical physics has been viewed as positive.

²Broadly speaking, theoretical physics [15]: motivates experiments and experimental facilities by proposing explanations for data that can be confirmed or falsified; interprets experimental results by connecting observations to abstract mathematical structures; suggests new experiments or analyses based on the predictions of models or theories; explores the consequences of theories and models in future experiments; extends the boundaries of our understanding of Nature and incorporates new perspectives (such as quantum information); develops new techniques relevant to experiments and builds tools used by the experimental community for simulation and analysis; and constructs theories and models that connect particle physics to gravity, cosmology, astrophysics, nuclear physics, condensed matter, atomic, molecular, and optical (AMO) physics, and mathematics.

However, the potential of AI-driven theory requires us to reconsider what is truly fundamental to theoretical physics.

Organizing at community scale would unlock assets that no individual researcher or research group can obtain:

- **Frontier-scale compute and model access:** resources beyond a single PI’s budget and not subject to the non-uniformity of university resources.
- **The agent-interaction corpus:** a collective record of how agents attack theory problems, including prompts, tool calls, generated code, failures, and human corrections which, when pooled across many groups, can be used to improve future agentic harnesses with proper human oversight to improve scientific alignment. Theorists have a direct scientific interest in the problem-solving strategies chosen by LLMs.
- **Shared validation infrastructure:** benchmarks for agentic theory workflows needed to build trust, and to establish which models suffice for which tasks. This is only useful if a wide variety of use cases are available [16–20].
- **Scientific foundation models:** training runs no single group can fund, whose value scales with the diversity of contributed data.

These assets did not exist five years ago, and there was no reason to, or mechanism for, the theory community to organize. With the demonstrated power of state-of-the-art AI and the launch of the GM, this situation has changed. To obtain these goods, the theory community needs to bind into a larger, organized and coherent structure.

Also, the current paradigm of “reinforcement learning with verifiable rewards” has led to rapid advances in agentic capabilities but also well-publicized incidents of agents exhibiting misaligned behaviors including cheating, unintended collusion with other agents, breaking out of sandboxes, illegal hacking, *etc.* One explanation is insufficient human oversight of powerful agents. In the theory domain, engaging with the broader theory community leverages human resources necessary to make sure theory agents remain properly aligned with scientific practices.

An analogous situation occurred over three decades ago in the rise of high performance computing (HPC) and the impact on a major consumer in theoretical HENP, lattice gauge theory. Technology was evolving rapidly with software co-evolving at the same pace and resources becoming obsolete every few years, beyond the capabilities of individual universities and small research groups to keep up. The U.S. lattice field theory community banded together to form the USQCD organization to coordinate efforts on keeping up with the latest advances in technology (hardware and software) while leaving researchers free to pursue their individual theoretical physics interests using the shared community technology. The mere fact that the organization has persisted and thrived for over two decades when other national theory-oriented collaborations have come and gone is evidence enough that it is a collaborative model for theorists worth emulating.

The organization of an HEP-experimental AI community (as outlined in [2]) is already underway. Theory should not duplicate that effort; we argue that theory AI research differs enough from its experimental counterpart to warrant its own organizational unit, even where the underlying infrastructure is best shared. A critical part of theory research in AI will be keeping the traditional, smaller groups connected to a larger effort **and** sharing what they learn, reporting both research successes as well as putative failures.

2 FOUR MODES OF AI IN PARTICLE THEORY

The following subsections detail the way theoretical physics can benefit from and contribute to the GM: agentic AI to accelerate standard theory tasks (Section 2.1); theory of and for AI to better understand why and how AI works and develop better applications (Section 2.2); bespoke AI for the

unique problems that particle theory aims to solve (Section 2.3); and the construction and study of scientific foundation models that could lead to a deeper understanding of the structure of nature (Section 2.4). The shared assets identified in Section 1.2 are critical to these developments: agentic AI is the principal consumer of frontier model access and the source of the agent-interaction corpus; the shared validation infrastructure can establish which models are trustworthy for which theory tasks. Frontier-scale computing resources, meanwhile, are needed to train scientific foundation models. Theorists not only benefit from AI but also contribute uniquely to its further development. The subject of what theory contributes is addressed in Section 2.2.

2.1 AGENTIC AI IN THEORETICAL PHYSICS

AI agents can take on different roles, from **supervised research assistant**, to **autonomous scientific assistant**, to **autonomous scientist(s)**, *etc.* Here, we outline how these different agents can accelerate particle theory and AI advances.

AI SUPERVISED RESEARCH ASSISTANT The AI supervised research assistant does nontrivial work in close interaction with its human operators, who provide guidance, feedback, and verification.

There are various early examples of theoretical physics results produced in this manner,³ such as the AI-guided discovery of an unexpected connection between different invariants of knots (objects foundational to certain QFTs [25]), producing a proven theorem [26], closed-form expressions for certain gluon amplitudes [27], and a resummation of large logarithms in collider event shapes [28].

Emerging tools are further expanding what is possible for theory-oriented AI. AI proof assistants can verify formal mathematics, helping theorists construct and check rigorous proofs interactively [29, 30]. LLMs can assist with direct symbolic calculation, reshaping conventional computational pipelines in formal theory [8, 9, 31, 32]. An “AI co-scientist” is becoming a practical reality, capable of generating hypotheses from theoretical data, which the human scientist evaluates, selects, and develops [4, 6, 33, 34].

AI AUTONOMOUS SCIENTIFIC ASSISTANT Much of our day-to-day work involves applying well-established methodology to carry out complex but predictable tasks requiring minimal novel insight to complete. This includes routine mathematical calculations, managing and running simulations, and basic data analysis [5, 10, 35–40].

Many of these tasks can be automated with current or near-term AI, including: (1) calculations across families of new physics models and associated experimental searches to set limits and optimize searches [9, 41–45]; (2) tests of different string theories and the exploration of the large parameter space of the string landscape [46–48]; (3) searches over quantum circuit designs, quantum algorithms, and related constructions [49–51]; (4) optimization of algorithmic hyperparameters; and (5) orchestration of computing campaigns [5, 52, 53].

AI AUTONOMOUS SCIENTIST(S) This use case envisions AI acting as autonomous scientists, or swarms of them, attacking nontrivial research problems [6, 34, 54, 55]. At its most powerful extreme, this may include closed-loop systems generating and exploring novel hypotheses. It is critical to investigate their viability, but realizing automated science workflows will require significant experimentation. Within this pursuit, there are two interrelated approaches: long-running autonomous “stateful” agents that maintain context and build on their own prior work over extended periods [6, 56]; and multi-agent orchestration, where multiple AI agents are deployed with complementary roles [13, 36, 44, 52, 57–60]. Both sit at the frontier of modern AI engineering, undergoing extensive exploration across science and technology. A natural testing ground for such systems is **formal theory**, where theoretical data arises from definitions and can be both generated

³Further examples in pure mathematics include solutions to some of the Erdős problems [21], the First Proof Challenge problems [22], and results in graph combinatorics [23, 24].

and validated without physical experiments [8, 32].

2.2 THEORY OF AND FOR AI

The relationship between theoretical physics and AI is bi-directional. AI can accelerate the study of physics, while the mathematical tools of physics can deepen our understanding of AI itself. Particle theorists have a long history of exploring and explaining complex mathematical structures. QFTs provide both good problems and good datasets for understanding how and why AI/ML works.

UNDERSTANDING AI/ML THROUGH PHYSICS At its core, ML is the science of constructing, optimizing, and using function approximations, most often via stochastic optimization. The mathematics underlying ML is statistics on function spaces, closely related to QFT and statistical mechanics, and the same mathematical tools developed for QFT can be applied to study AI/ML [61, 62]. Physical reasoning can also shed light on learning dynamics and improve interpretability.

Neural network quantum field theory (NNQFT) [63, 64] sits at the intersection of understanding and advancing both AI and physics. NNQFT provides constructive definitions of QFTs, offering new insight and a fundamentally new approach to QFT itself [65].

A powerful application of QFT tools to AI/ML is for uncertainty quantification, an open challenge of critical importance across the sciences. In this context, QFT tools allow us to characterize the priors and biases encoded by different ML setups.

ADVANCING AI/ML WITH PHYSICS Beyond understanding existing methods, there is the potential to advance AI/ML with new physics-inspired frameworks and architectures. Diffusion models, the state-of-the-art for image generation and an emerging approach to language modeling, are defined in terms of stochastic processes, the same mathematics behind Brownian motion [66–68]. Innovations in the theory of ML, such as the development of symmetry-equivariant architectures driven by theoretical work in lattice field theory and collider phenomenology, give measurable gains in lattice gauge theory sampling and may have transformative impact on other scientific computations [69–73].

2.3 BESPOKE AI/ML FOR PHYSICS

High-quality learned approximations in theoretical physics, derived with custom AI/ML, can be useful for learning structure from “theoretical” and experimental data, accelerating simulations and sampling, and constructing theory-informed models for enhanced sensitivity or inference. Agentic AI can itself be used to accelerate progress in this area [14, 32, 74].

LEARNING FROM THEORETICAL AND PHYSICS DATA AI/ML can help explore new facets of pure mathematics and formal theoretical physics, including QFT and string theory. This exploration is characterized by manipulations of symbolic formulae or “theoretical data”: data arising from definitions that is obtained via mathematical classification or algorithmic generation. Examples include fundamental objects of abstract algebra, topology, and differential geometry: groups and their representations, topological spaces and their invariants, and metrics with special characteristics. Such data is of fundamental mathematical interest, but also of physical interest due to their role in defining physics theories or their associated observables.

Early, concrete achievements include the use of neural networks to uncover hidden mathematical relations in string theory data [75–78], driving conjecture generation via interpretability that led to rigorous new theorems. Further opportunities exist in computing Calabi-Yau metrics, Kaluza-Klein modes, and Yukawa couplings via learned methods [79–81], as well as in applying ML to QFT amplitudes [31, 32]. More broadly, symbolic regression, data mining, and related techniques offer promising avenues for extracting useful patterns from physics data and discovering new functional relationships [48, 82–84].

ACCELERATING SIMULATIONS AND SAMPLING Particle physics benefits from sophisticated numerical simulations of particle collisions, lattice field theories, cosmology, and more; in practice, discoveries, measurements, and limits are based on comparisons of simulator predictions to data. AI/ML methods can address the significant computational bottlenecks that limit the use and further development of these simulation tools [85–87]. Progress has already been seen in: (1) improving speed and efficiency over traditional Monte Carlo methods [69, 88–90]; (2) integrating parameter differentiability to enable fitting and uncertainty estimation [91]; (3) enabling on-the-fly inference [92, 93]; and (4) providing alternative physics models for less well-known physics [94–98].

THEORY-INFORMED ML MODELS Task-specific ML models can incorporate physics constraints and assumptions (symmetries, equivariances, new physics theory scenarios) directly informed by theory. These may enhance signal sensitivity in searches beyond what generic methods achieve.

An important application of AI is solving large-scale inverse problems. Particle phenomenology both supplies data and is reciprocally guided by many inverse problems, including effective field theories [99–101], UV-complete theory parameterizations [9, 55, 102], parton distribution functions [103], and hadronization models [94, 96, 97]. AI/ML can be closely integrated into the parametrization and inference tasks in this space, enabling new approaches to uncertainty quantification, validation, and verification with consistent treatments of theoretical and experimental uncertainties carried through the full phenomenology workflow [92, 104, 105].

AI/ML FOR QUANTUM INFORMATION SCIENCE AI can help mitigate the theoretical complexity of QFT and the physical limitations of current quantum hardware [106]. Driven by the bottleneck of classical computing, a concerted effort is underway to build the theoretical frameworks and algorithms necessary to simulate QFT on both digital quantum computers and analog simulators. Custom AI/ML methods have the potential to accelerate this effort across multiple fronts, from quantum circuit optimization and algorithm construction to learned error mitigation strategies [49–51]. The community needs to converge on a set of domain-specific benchmarks that can demonstrate empirical quantum advantage on QFT problems.

2.4 SCIENTIFIC FOUNDATION MODELS

Recent progress in AI has shown that large, flexible parameterizations trained to simultaneously model large datasets across different modalities (text, image, video, audio) can develop general, abstract representations of the same object in different modalities. This manifests as the ability to describe and reason about image and video data, or to produce images and video from language inputs. Applying this same functionality to physics data from experiments, observations, and simulations could yield powerful “scientific foundation models” [107–112] whose impact extends well beyond the custom, task-specific methods described above.

Beyond their direct utility, scientific foundation models are themselves useful objects of study, since they encode, in some fashion, a theory describing the physics data. Interpretability and explainability techniques [113] and learned representations [114], *etc.*, may reveal characterizations of these implicit theories that are comprehensible by humans. It may be possible to derive useful effective theories, achieve more detailed characterizations of experimental artifacts, or discover new physics hidden in data that was missed by traditional analysis techniques.

More speculatively, physics data could be incorporated into a generalized vision-language model [115]. Presently, both humans and AI are constrained to clever data reductions to map high-dimensional, heterogeneous physics data into tractable summaries [116]. However, an AI model could natively ingest and reason about such data without lossy intermediate reductions, exposing powerful new capabilities and improving our ability to identify interesting but subtle features in complex physics datasets. From the AI perspective, particle physics provides an unexploited reservoir of highly structured data to help improve AI’s general reasoning capabilities.

3 BUILDING AN AI THEORY COMMUNITY FOR THE GENESIS MISSION

We have argued that a large-scale collaboration among theorists, similar to the organization in frontier experiments, provides many shared assets: frontier-scale computing resources and access to frontier models, the large amounts of data that constitute the agent-interaction corpus, shared validation infrastructure for benchmarking, and scientific foundation models. These shared assets are particularly valuable to theorists. Individual theory groups typically do not command compute or model budgets on the scale that experimental collaborations pool collectively (Section 1.2). While theorists can share infrastructure with an experiment-oriented effort, there are special considerations for AI-driven theory. Theorists work differently with AI, often in long, exploratory, interactive sessions without a guaranteed outcome. This methodology does not fit well with batch-computing and allocations based on promised outcomes, a model that governs most existing computing infrastructure (Section 3.4). We have further outlined the type of research theorists are doing and can do with AI. While there is overlap with experimental research, theory is often unique and valuable to particle physics as a whole. To empower this research, theory needs to have a voice in the organization of a field-wide AI effort, and that requires a community.

The national laboratory theory groups can serve as a backbone for building the larger collaborative units that we need to accomplish the type of vision presented here. By employing staff scientists and hosting and maintaining shared services such as the agent-interaction repository, the national laboratories can provide the institutional continuity that individual university groups cannot easily sustain against the mobility of graduate students and postdoctoral researchers.

We provide concrete steps for forming and developing this community in Section 4. Realizing this vision will require attention to access, infrastructure, policy, and education. As the Genesis Mission Platform is made available, we expect that large, well-supported groups centered at national laboratories will have the resources available to efficiently leverage it. Lack of resources may be an entry barrier for individual researchers at universities, particularly those who are less interested in AI infrastructure and more interested in research problems. This pattern is consistent with past experience: during the development of the Leadership Computing Facilities (LCFs) at the national labs, USQCD played an important coordination role, enabling access for individual researchers to LCF resources and providing assurances, through the use of the USQCD Software API, that valuable LCF resources would be used efficiently. The proposed community organization will provide analogous services with the Genesis Mission Platform.

3.1 PROCUREMENT AND MODEL ACCESS

There is no substitute for general, fluid intelligence in theoretical physics. Theoretical physics proceeds by solving challenging problems, each solution exposing new problems in turn. To accelerate this process with AI, general quantitative problem-solving ability is necessary, and there is accordingly no substitute for access to the best models available at the frontier of AI capabilities. This has important implications for procurement.

Model capabilities are improving rapidly, and multiple major frontier providers are competing to develop the best AI models. A snapshot of the present ranking of model capabilities is unlikely to extrapolate reliably beyond even a few months. It is important that procurement decisions are flexible enough to allow access to the best models at any given moment without locking into a particular model family or provider.

There are broad efforts to measure the overall intelligence and utility of AI models by constructing benchmarks [16, 17, 19, 20, 56, 117]. These are important and useful, but also highly limited and should not be over interpreted. There are well-known issues with the useful lifetime of any benchmark, which is limited by both saturation as models improve and “benchmaxing,” or training to the test, which furthermore complicates their interpretation. More fundamentally, a single scalar

metric cannot usefully capture capabilities across the broad skillset required for theoretical physics. While it will be useful to further develop theory-specific benchmarks, procurement decisions should not be made solely on these grounds. Instead, a broad range of models should be made available to theorists, with an aggregation of their opinions as the primary deciding factor on utility. Benchmarks, nonetheless, inform these decisions in a different way, by roughly establishing which class of models reliably automate a given task and therefore where frontier access may not be required.

3.2 OPEN-WEIGHTS MODELS

Open-weights (ideally, open-source) models are obviously desirable from an academic perspective. From a procurement perspective, they can also be offered at significantly reduced cost compared to frontier models. Open-weights models are increasingly designed for efficiency. Moreover, they can be hosted locally, on either institutional hardware or the cloud, providing bare-minimum prices modulo economies of scale. Hardware purchased to serve these models can also be used for other computational science tasks, providing additional benefits. In principle, locally hosted models also allow more reliable guarantees of information security: data is never streamed to any corporate entity.

However, at time of writing, open models are not an equivalent substitute for frontier models and in practice are qualitatively less useful for general theoretical physics tasks. While open models developed by industrial and academic efforts [118, 119] are improving at similar rates to closed frontier models, there is a well-known lag. Common heuristics put it in the range of six months to a year. Further investment in open model development may close this gap. However, if institutional decisions force theorists to rely solely on these options, they will be disadvantaged.

3.3 MODES OF ACCESS

Virtually everyone with an Internet connection has access to some form of chatbot AI. However, as already emphasized, access to the best available models is critical for theory needs. This is especially true in the highly interactive chatbot mode, which delivers outputs directly to users; low-quality outputs result in confusion and wasted human time. At present price levels, access to frontier models is affordable but nontrivial for individual users. For junior career researchers, it can be a substantial financial burden. Policies should be adjusted to fund this access. Information security is also a concern.

Access to AI in the now-familiar “chatbot” form is important, but only exposes a small fraction of the capabilities of emerging agentic use-cases. Off-the-shelf proprietary agentic tools like Claude Code, Codex, and similar products are proving to be extremely useful science tools [11, 12, 28]. Enabling custom-built agentic AI requires API-level access [52, 120, 121], which can be substantially more expensive than subscription-based access. More aggressive uses of agentic AI cannot be expected to be paid for out-of-pocket by researchers and should be funded by grants and institutions.

3.4 HARDWARE AND COMPUTING INFRASTRUCTURE

Using agentic AI imposes hardware needs distinct from “inference service” hardware where the actual language model is run. Modern agent systems are not fully web-native in their most useful mode. Instead, they read and write files, interact with other data sources, and, critically, run programs and execute code. This requires a computing environment for the agents to work within [122].

Much can be enabled simply with high-quality laptops or workstations. For junior-career researchers, this can be a significant financial burden. At the policy level, it will be important for grants to account for this in equipment budgets. At the group level, PIs should be encouraged to provide this now-critical equipment to their graduate students and postdocs.

However, there are limitations to what can be accomplished on a laptop, especially in computa-

tionally intensive pursuits. There is a critical and unfilled need to provide broad and easy access to larger-scale computing resources to allow experimentation. A combination of group and institutional clusters and cloud computing provide the simplest solutions, but may be financially prohibitive, and moreover have an upper bound on scale.

At present, allocations are granted at LCFs and other supercomputing facilities based on promised outcomes. This is well-motivated for traditional computational science, but incompatible with enabling AI/ML research. AI/ML research generically requires computational experiments with no guaranteed physics output: the experiment itself, and how well it performed, is the useful output. This includes both training of new models, as well as automated computational science with agentic AI. Most of this type of R&D is “wasted compute” in the framework of traditional allocation structures, but is critical to make progress. This is the role that can and should be filled by the American Science Cloud (AmSC), but should also be addressed by updating allocation policies. Efforts in this direction are already underway: the USQCD Collaboration, for example, is currently engaged with the AmSC as an infrastructure partner and is supporting seed AI model teams.

3.5 EDUCATION AND WORKFORCE DEVELOPMENT

Realizing the vision described in this document requires that theorists at all career stages develop fluency with AI tools. This represents a significant shift in the skillset expected of theoretical physicists. Not every theorist must become an AI/ML expert, but a working familiarity with these tools (their capabilities and limitations) should become standard.

At the graduate level, training should focus on using AI tools for physics research and understanding the principles behind them.

For established researchers, the barrier is different: many are skeptical of or unfamiliar with what AI tools can do. This can be addressed with accessible demonstrations, workshops, and community forums that showcase practical applications relevant to their specific research areas.

It is especially important to support junior-career researchers (graduate students and postdocs) who may be the most natural early adopters but who also face the greatest financial and institutional barriers to access.

Mentoring norms, on-boarding resources, career guidance and skill development, recognition and evaluation criteria, and other issues need to be discussed in an open forum.

4 CONCLUSIONS

The use of AI in the natural sciences is a reality and an opportunity to advance HEP theory while contributing to the general AI eco-system. However, the exploitation of this opportunity requires structural change in the theory community. Theorists can still maintain the value of their unique practice while finding ways to share resources. The USQCD model has demonstrated this can work.

Concretely, we propose the following near-term sequence. First, national laboratory theory groups convene an initial community meeting, within the coming year, to establish the scope of the organization: how frontier model access is pooled and procured, what shared benchmarks the community most needs, how agentic interaction data is formatted and governed, and what a community-scale foundation-model effort would require. Second, the participating labs create the first of these shared services, a pilot agent-interaction repository on the Genesis Mission Platform, open to any contributing theory group regardless of institutional affiliation. We propose this service first, because it is the least expensive one to build, and its value grows fastest with the number of contributing groups. Third, these services are connected to the broader national effort through the Genesis Mission Platform, interoperating with the experimental community’s corresponding infrastructure [2]

through common formats.

REFERENCES

- [1] *Genesis Mission National Science and Technology Challenges*. 2026. URL: <https://www.energy.gov/documents/genesis-mission-science-and-technology-challenges>.
- [2] Thea Klæboe Aarrestad et al. *Building an AI-native Research Ecosystem for Experimental Particle Physics: A Community Vision*. Feb. 2026. arXiv: 2602.17582 [hep-ex].
- [3] Junjie Li et al. *Agent Harness Engineering: A Survey*. 2026. URL: <https://openreview.net/pdf?id=e0Nq7FdiHa>.
- [4] S. D. Bakshi et al. “ArgoLOOM: agentic AI for fundamental physics from quarks to cosmos”. In: (Oct. 2025). arXiv: 2510.02426 [hep-ph].
- [5] Tony Menzo et al. “HEPTAPOD: Orchestrating High Energy Physics Workflows Towards Autonomous Agency”. In: (Dec. 2025). arXiv: 2512.15867 [hep-ph].
- [6] Tingjia Miao et al. “PhysMaster: Building an Autonomous AI Physicist for Theoretical and Computational Physics Research”. In: (Dec. 2025). arXiv: 2512.19799 [cs.AI].
- [7] Prateek Agrawal et al. “The FERMIACC: Agents for Particle Theory”. In: (Mar. 2026). arXiv: 2603.22538 [hep-ph].
- [8] Tony Menzo et al. “Agentic Diagrammatica: Towards Autonomous Symbolic Computation in High Energy Physics”. In: (Mar. 2026). arXiv: 2603.26990 [hep-ph].
- [9] Shaikh Saad. “Large Language Model-Assisted Framework for BSM Model Building”. In: (June 2026). arXiv: 2606.21316 [hep-ph].
- [10] Eli Gendreau-Distler et al. “Automating High Energy Physics Data Analysis with LLM-Powered Agents”. In: *39th Annual Conference on Neural Information Processing Systems: Includes Machine Learning and the Physical Sciences (ML4PS)*. Dec. 2025. arXiv: 2512.07785 [physics.data-an].
- [11] Eric A. Moreno et al. “AI Agents Can Already Autonomously Perform Experimental High Energy Physics”. In: (Mar. 2026). arXiv: 2603.20179 [hep-ex].
- [12] Anthony Badea et al. “Agentic AI – Physicist Collaboration in Experimental Particle Physics: A Proof-of-Concept Measurement with LEP Open Data”. In: (Mar. 2026). arXiv: 2603.05735 [hep-ex].
- [13] Mingfeng He et al. “Dr.Sai: An agentic AI for real-world physics analysis at BESIII”. In: (Apr. 2026). arXiv: 2604.22541 [hep-ex].
- [14] Sascha Diefenbacher et al. “Agents of Discovery”. In: (Sept. 2025). arXiv: 2509.08535 [hep-ph].
- [15] N. Craig et al. *Snowmass Theory Frontier Report*. Nov. 2022. arXiv: 2211.05772 [hep-ph].
- [16] Daniel J. H. Chung et al. “Theoretical physics benchmark (TPBench)—a dataset and study of AI reasoning capabilities in theoretical physics”. In: *Mach. Learn. Sci. Tech.* 6.3 (2025), p. 030505. DOI: 10.1088/2632-2153/adfcb0. arXiv: 2502.15815 [cs.LG].
- [17] Minhui Zhu et al. “Probing the Critical Point (CritPt) of AI Reasoning: a Frontier Physics Research Benchmark”. In: (Sept. 2025). arXiv: 2509.26574 [cs.AI].
- [18] Ben Hawks et al. “An MLCommons Scientific Benchmarks Ontology”. In: (Nov. 2025). arXiv: 2511.05614 [cs.LG].
- [19] Darius A. Faroughy et al. “Collider-Bench: Benchmarking AI Agents with Particle Physics Analysis Reproduction”. In: (May 2026). arXiv: 2605.13950 [cs.LG].

- [20] Tingjia Miao et al. “PRL-Bench: A Comprehensive Benchmark Evaluating LLMs’ Capabilities in Frontier Physics Research”. In: (Apr. 2026). arXiv: 2604.15411 [cs.LG].
- [21] Tony Feng et al. *Semi-Autonomous Mathematics Discovery with Gemini: A Case Study on the Erdős Problems*. Jan. 2026. arXiv: 2601.22401 [math.CO].
- [22] Mohammed Abouzaid et al. *First Proof Second Batch*. June 2026. arXiv: 2606.18119 [math.HO].
- [23] Adam Zsolt Wagner. *Constructions in combinatorics via neural networks*. 2021. arXiv: 2104.14516 [math.CO].
- [24] Alexander Novikov et al. *AlphaEvolve: A coding agent for scientific and algorithmic discovery*. 2025. arXiv: 2506.13131 [cs.AI].
- [25] Edward Witten. “Quantum Field Theory and the Jones Polynomial”. In: *Commun. Math. Phys.* 121 (1989), pp. 351–399.
- [26] Alex Davies et al. “Advancing mathematics by guiding human intuition with AI”. In: *Nature* 600 (2021), pp. 70–74.
- [27] Alfredo Guevara et al. “Single-minus gluon tree amplitudes are nonzero”. In: (Feb. 2026). arXiv: 2602.12176 [hep-th].
- [28] Matthew D. Schwartz. “Resummation of the C-Parameter Sudakov Shoulder Using Effective Field Theory”. In: (Jan. 2026). arXiv: 2601.02484 [hep-ph].
- [29] The mathlib Community. *The Lean mathematical library*. 2020. DOI: 10.1145/3372885.3373824. arXiv: 1910.09336 [cs.LO].
- [30] Thomas Hubert et al. “Olympiad-level formal mathematical reasoning with reinforcement learning”. In: *Nature* 651.8106 (2026), pp. 607–613. DOI: 10.1038/s41586-025-09833-y.
- [31] Abdulhakim Alnuqaydan, Sergei Gleyzer, and Harrison Prosper. “SYMBA: symbolic computation of squared amplitudes in high energy physics with machine learning”. In: *Mach. Learn. Sci. Tech.* 4.1 (2023), p. 015007. DOI: 10.1088/2632-2153/acb2b2. arXiv: 2206.08901 [hep-ph].
- [32] Yi Gu and Sven Krippendorff. “Scattering Amplitudes as Programs: Self-Evolving Search for Theory and Event Generation”. In: (July 2026). arXiv: 2607.21629 [hep-ph].
- [33] Juraj Gottweis et al. *Accelerating scientific discovery with Co-Scientist*. 2025. arXiv: 2502.18864 [cs.AI].
- [34] Chris Lu et al. *The AI Scientist: Towards Fully Automated Open-Ended Scientific Discovery*. 2024. arXiv: 2408.06292 [cs.AI].
- [35] Shi Qiu et al. “An End-to-end Architecture for Collider Physics and Beyond”. In: (Mar. 2026). arXiv: 2603.14553 [hep-ph].
- [36] Tilman Plehn, Daniel Schiller, and Nikita Schmal. “MadAgents”. In: (Jan. 2026). arXiv: 2601.21015 [hep-ph].
- [37] Aman Desai. “RooAgent: An LLM Agent for Root-Based High Energy Physics Analysis”. In: (May 2026). arXiv: 2605.17318 [hep-ph].
- [38] Junkun Jiao et al. “HepScript: A Dual-Use DSL for Human-AI Collaborative Data Analysis Workflows in High-Energy Physics”. In: (May 2026). arXiv: 2605.01423 [hep-ex].
- [39] Haofer Gao et al. “LQCDMaster: Agentic Scientific Computing for Lattice Quantum Chromodynamics Research”. In: (July 2026). arXiv: 2607.15001 [hep-lat].
- [40] Qi Wang et al. “CLVisc Agent for autonomous relativistic hydrodynamics studies”. In: (July 2026). arXiv: 2607.27822 [nucl-th].

- [41] Yang Xiao, Yuanfang Yue, and Yang Zhang. “EasyScan_HEP 2: Agent-Ready Parameter Scans for High-Energy Physics”. In: (June 2026). arXiv: 2606.31214 [hep-ph].
- [42] Yu-Chen Guo, Jie Wang, and Ji-Chong Yang. “SMEFT-Pheno-Agent: a natural-language-driven AI agent for machine-learning-assisted Standard Model Effective Field Theory phenomenology”. In: (July 2026). arXiv: 2607.22331 [hep-ph].
- [43] Sascha Diefenbacher et al. “Agentic Re-Casting using Agentic Re-Simulations”. In: (July 2026). arXiv: 2607.22813 [hep-ph].
- [44] Michele Lucente et al. “DarkAgents”. In: (June 2026). arXiv: 2606.11157 [hep-ph].
- [45] Isaac R. Wang. “LeWRON: Agentic Analysis of Electroweak Phase Transitions”. In: (June 2026). arXiv: 2606.19425 [hep-ph].
- [46] Fabian Ruehle. “Data science applications to string theory”. In: *Phys. Rept.* 839 (2020), pp. 1–117. DOI: 10.1016/j.physrep.2019.09.005.
- [47] James Halverson, Brent Nelson, and Fabian Ruehle. “Branes with Brains: Exploring String Vacua with Deep Reinforcement Learning”. In: *JHEP* 06 (2019), p. 003. DOI: 10.1007/JHEP06(2019)003. arXiv: 1903.11616 [hep-th].
- [48] Konstantin T. Matchev et al. “Exploring the truth and beauty of theory landscapes with machine learning”. In: *Phys. Lett. B* 856 (2024), p. 138941. DOI: 10.1016/j.physletb.2024.138941. arXiv: 2401.11513 [hep-ph].
- [49] Francisco J. R. Ruiz et al. “Quantum Circuit Optimization with AlphaTensor”. In: (Feb. 2024). arXiv: 2402.14396 [quant-ph].
- [50] Qing-Hong Cao et al. “Scalable Quantum State Preparation via Large-Language-Model-Driven Discovery”. In: (May 2025). arXiv: 2505.06347 [quant-ph].
- [51] Marco Knipfer et al. “AI Agents for Variational Quantum Circuit Design”. In: (Feb. 2026). arXiv: 2602.19387 [quant-ph].
- [52] Alexander Roman and Jacob Roman. “Orchestral AI: A Framework for Agent Orchestration”. In: (Jan. 2026). arXiv: 2601.02577 [cs.AI].
- [53] Pietro Lugato et al. “Archi: Agentic Operations at the CMS Experiment”. In: (June 2026). arXiv: 2606.04755 [hep-ex].
- [54] Jin-Xin Tan et al. “Automated Extraction of Collins-Soper Kernel from Lattice QCD using An Autonomous AI Physicist System”. In: (Mar. 2026). arXiv: 2603.22471 [hep-lat].
- [55] Stephon Alexander et al. “Autonomous Discovery of Particle Physics Theories from Experimental Data”. In: (Mar. 2026). arXiv: 2603.28935 [hep-ph].
- [56] Thomas Kwa, Ben West, Joel Becker, et al. *Measuring AI Ability to Complete Long Software Tasks*. 2025. arXiv: 2503.14499 [cs.AI].
- [57] Joschka Birk et al. “A Scientific Human-Agent Reproduction Pipeline”. In: (Apr. 2026). DOI: 10.5281/zenodo.21078068. arXiv: 2604.18752 [hep-ph].
- [58] Antonio J. Costa et al. “AgentRivet: an automated system for producing Rivet routines from journal publications”. In: (June 2026). arXiv: 2606.13535 [hep-ex].
- [59] Ahmed Hammad and Mihoko Nojiri. “Articulating Assumptions in AI-Generated Scientific Analyses through Task Decomposition”. In: (July 2026). arXiv: 2607.05762 [cs.SE].
- [60] Masahiko Saito, Tomoe Kishimoto, and Junichi Tanaka. “Development of an LLM-Based System for Automatic Code Generation from HEP Publications”. In: Apr. 2026. arXiv: 2604.14696 [physics.data-an].

- [61] Daniel A. Roberts, Sho Yaida, and Boris Hanin. *The Principles of Deep Learning Theory*. Cambridge University Press, May 2022. ISBN: 978-1-009-02340-5. DOI: 10.1017/9781009023405. arXiv: 2106.10165 [cs.LG].
- [62] Jim Halverson. “TASI Lectures on Physics for Machine Learning”. In: (July 2024). arXiv: 2408.00082 [hep-th].
- [63] James Halverson, Anindita Maiti, and Keegan Stoner. “Neural Networks and Quantum Field Theory”. In: *Mach. Learn. Sci. Tech.* 2.3 (2021), p. 035002. DOI: 10.1088/2632-2153/abeca3. arXiv: 2008.08601 [cs.LG].
- [64] Mehmet Demirtas et al. “Neural network field theories: non-Gaussianity, actions, and locality”. In: *Mach. Learn. Sci. Tech.* 5.1 (2024), p. 015002. DOI: 10.1088/2632-2153/ad17d3. arXiv: 2307.03223 [hep-th].
- [65] Christian Ferko, James Halverson, and Aaron Mutchler. “Universality of Neural Network Field Theory”. In: (Jan. 2026). arXiv: 2601.14453 [hep-th].
- [66] Jascha Sohl-Dickstein et al. “Deep Unsupervised Learning using Nonequilibrium Thermodynamics”. In: (Mar. 2015). arXiv: 1503.03585 [cs.LG].
- [67] Jonathan Ho, Ajay Jain, and Pieter Abbeel. “Denoising Diffusion Probabilistic Models”. In: (June 2020). arXiv: 2006.11239 [cs.LG].
- [68] Shen Nie et al. *Large Language Diffusion Models*. 2025. arXiv: 2502.09992 [cs.CL].
- [69] M. S. Albergo, G. Kanwar, and P. E. Shanahan. “Flow-based generative models for Markov chain Monte Carlo in lattice field theory”. In: *Phys. Rev. D* 100.3 (2019), p. 034515. DOI: 10.1103/PhysRevD.100.034515. arXiv: 1904.12072 [hep-lat].
- [70] Gurtej Kanwar et al. “Equivariant flow-based sampling for lattice gauge theory”. In: *Phys. Rev. Lett.* 125.12 (2020), p. 121601. DOI: 10.1103/PhysRevLett.125.121601. arXiv: 2003.06413 [hep-lat].
- [71] Denis Boyda et al. “Sampling using $SU(N)$ gauge equivariant flows”. In: *Phys. Rev. D* 103.7 (2021), p. 074504. DOI: 10.1103/PhysRevD.103.074504. arXiv: 2008.05456 [hep-lat].
- [72] Shiqi Gong et al. “An efficient Lorentz equivariant graph neural network for jet tagging”. In: *JHEP* 07 (2022), p. 030. DOI: 10.1007/JHEP07(2022)030. arXiv: 2201.08187 [hep-ph].
- [73] Huilin Qu, Congqiao Li, and Sitian Qian. “Particle Transformer for Jet Tagging”. In: (Feb. 2022). arXiv: 2202.03772 [hep-ph].
- [74] W. Esmail, A. Hammad, and M. Nojiri. “CoLLM: AI engineering toolbox for end-to-end deep learning in collider analyses”. In: (Feb. 2026). arXiv: 2602.06496 [hep-ph].
- [75] Jonathan Carifio et al. “Machine Learning in the String Landscape”. In: *JHEP* 09 (2017), p. 157. arXiv: 1707.00655 [hep-th].
- [76] Yang-Hui He. *Deep-Learning the Landscape*. 2017. arXiv: 1706.02714 [hep-th].
- [77] Daniel Krefl and Rak-Kyeong Seong. “Machine Learning of Calabi-Yau Volumes”. In: *Phys. Rev. D* 96 (2017), p. 066014. arXiv: 1706.03346 [hep-th].
- [78] Fabian Ruehle. “Evolving neural networks with genetic algorithms to study the String Landscape”. In: *JHEP* 08 (2017), p. 038. arXiv: 1706.07024 [hep-th].
- [79] Lara B. Anderson et al. “Moduli-dependent Calabi-Yau and $SU(3)$ -structure metrics from Machine Learning”. In: *JHEP* 05 (2021), p. 013. DOI: 10.1007/JHEP05(2021)013. arXiv: 2012.04656 [hep-th].
- [80] Michael R. Douglas, Subramanian Lakshminarasimhan, and Yidi Qi. “Numerical Calabi-Yau metrics from holomorphic networks”. In: (Dec. 2020). arXiv: 2012.04797 [hep-th].

- [81] Vishnu Jejjala, Damian Kaloni Mayorga Pena, and Challenger Mishra. “Neural network approximations for Calabi-Yau metrics”. In: *JHEP* 08 (2022), p. 105. DOI: 10.1007/JHEP08(2022)105. arXiv: 2012.15821 [hep-th].
- [82] Silviu-Marian Udrescu and Max Tegmark. “AI Feynman: a Physics-Inspired Method for Symbolic Regression”. In: *Sci. Adv.* 6.16 (2020), eaay2631. DOI: 10.1126/sciadv.aay2631. arXiv: 1905.11481 [physics.comp-ph].
- [83] Zhongtian Dong et al. “Is the machine smarter than the theorist: Deriving formulas for particle kinematics with symbolic regression”. In: *Phys. Rev. D* 107.5 (2023), p. 055018. DOI: 10.1103/PhysRevD.107.055018. arXiv: 2211.08420 [hep-ph].
- [84] Miles Cranmer. *Interpretable Machine Learning for Science with PySR and SymbolicRegression.jl*. 2023. arXiv: 2305.01582 [astro-ph.IM].
- [85] Matthew Feickert and Benjamin Nachman. “A Living Review of Machine Learning for Particle Physics”. In: (Feb. 2021). arXiv: 2102.02770 [hep-ph].
- [86] Phiala Shanahan et al. “Snowmass 2021 Computational Frontier CompF03 Topical Group Report: Machine Learning”. In: (Sept. 2022). arXiv: 2209.07559 [physics.comp-ph].
- [87] Simon Badger et al. “Machine learning and LHC event generation”. In: *SciPost Phys.* 14.4 (2023). Ed. by Anja Butter, Tilman Plehn, and Steffen Schumann, p. 079. DOI: 10.21468/SciPostPhys.14.4.079. arXiv: 2203.07460 [hep-ph].
- [88] Enrico Bothmann et al. “Exploring phase space with Neural Importance Sampling”. In: *SciPost Phys.* 8.4 (2020), p. 069. DOI: 10.21468/SciPostPhys.8.4.069. arXiv: 2001.05478 [hep-ph].
- [89] Christina Gao et al. “Event Generation with Normalizing Flows”. In: *Phys. Rev. D* 101.7 (2020), p. 076002. DOI: 10.1103/PhysRevD.101.076002. arXiv: 2001.10028 [hep-ph].
- [90] Konstantin T. Matchev and Prasanth Shyamsundar. “OASIS: Optimal Analysis-Specific Importance Sampling for event generation”. In: *SciPost Phys.* 10.2 (2021), p. 034. DOI: 10.21468/SciPostPhys.10.2.034. arXiv: 2006.16972 [hep-ph].
- [91] Lukas Heinrich and Michael Kagan. “Differentiable Matrix Elements with MadJax”. In: *J. Phys. Conf. Ser.* 2438.1 (2023), p. 012137. DOI: 10.1088/1742-6596/2438/1/012137. arXiv: 2203.00057 [hep-ph].
- [92] Kyle Cranmer, Johann Brehmer, and Gilles Louppe. “The frontier of simulation-based inference”. In: *Proc. Nat. Acad. Sci.* 117.48 (2020), pp. 30055–30062. DOI: 10.1073/pnas.1912789117. arXiv: 1911.01429 [stat.ML].
- [93] Kyoungchul Kong et al. “New Machine Learning Techniques for Simulation-Based Inference: InferenceStatic Nets, Kernel Score Estimation, and Kernel Likelihood Ratio Estimation”. In: (Oct. 2022). arXiv: 2210.01680 [stat.ML].
- [94] Phil Ilten et al. “Modeling hadronization using machine learning”. In: *SciPost Phys.* 14.3 (2023), p. 027. DOI: 10.21468/SciPostPhys.14.3.027. arXiv: 2203.04983 [hep-ph].
- [95] Jay Chan et al. “Fitting a deep generative hadronization model”. In: *JHEP* 09 (2023), p. 084. DOI: 10.1007/JHEP09(2023)084. arXiv: 2305.17169 [hep-ph].
- [96] Christian Bierlich et al. “Towards a data-driven model of hadronization using normalizing flows”. In: *SciPost Phys.* 17.2 (2024), p. 045. DOI: 10.21468/SciPostPhys.17.2.045. arXiv: 2311.09296 [hep-ph].
- [97] Benoît Assi et al. “Characterizing the hadronization of parton showers using the HOMER method”. In: *SciPost Phys.* 19.5 (2025), p. 125. DOI: 10.21468/SciPostPhys.19.5.125. arXiv: 2503.05667 [hep-ph].

- [98] Anja Butter et al. “Iterative HOMER with uncertainties”. In: *SciPost Phys.* 20.2 (2026), p. 042. DOI: 10.21468/SciPostPhys.20.2.042. arXiv: 2509.03592 [hep-ph].
- [99] Johann Brehmer et al. “Constraining Effective Field Theories with Machine Learning”. In: *Phys. Rev. Lett.* 121.11 (2018), p. 111801. DOI: 10.1103/PhysRevLett.121.111801. arXiv: 1805.00013 [hep-ph].
- [100] Jacob J. Ethier et al. “Combined SMEFT interpretation of Higgs, diboson, and top quark data from the LHC”. In: *JHEP* 11 (2021), p. 089. DOI: 10.1007/JHEP11(2021)089. arXiv: 2105.00006 [hep-ph].
- [101] Raquel Gomez Ambrosio et al. “Unbinned multivariate observables for global SMEFT analyses from machine learning”. In: *JHEP* 03 (2023), p. 033. DOI: 10.1007/JHEP03(2023)033. arXiv: 2211.02058 [hep-ph].
- [102] T. R. Harvey and A. Lukas. “Quark Mass Models and Reinforcement Learning”. In: *JHEP* 08 (2021), p. 161. DOI: 10.1007/JHEP08(2021)161. arXiv: 2103.04759 [hep-th].
- [103] Richard D. Ball et al. “The path to proton structure at 1% accuracy”. In: *Eur. Phys. J. C* 82.5 (2022), p. 428. DOI: 10.1140/epjc/s10052-022-10328-7. arXiv: 2109.02653 [hep-ph].
- [104] Johann Brehmer et al. “MadMiner: Machine learning-based inference for particle physics”. In: *Comput. Softw. Big Sci.* 4.1 (2020), p. 3. DOI: 10.1007/s41781-020-0035-2. arXiv: 1907.10621 [hep-ph].
- [105] Johann Brehmer and Kyle Cranmer. “Simulation-based inference methods for particle physics”. In: (Oct. 2020). arXiv: 2010.06439 [hep-ph].
- [106] Simon Catterall et al. “Report of the Snowmass 2021 Theory Frontier Topical Group on Quantum Information Science”. In: *Snowmass 2021*. Sept. 2022. DOI: 10.2172/1892238. arXiv: 2209.14839 [quant-ph].
- [107] Tobias Golling et al. “Masked particle modeling on sets: towards self-supervised high energy physics foundation models”. In: *Mach. Learn. Sci. Tech.* 5.3 (2024), p. 035074. DOI: 10.1088/2632-2153/ad64a8. arXiv: 2401.13537 [hep-ph].
- [108] Joschka Birk, Anna Hallin, and Gregor Kasieczka. “OmniJet- α : the first cross-task foundation model for particle physics”. In: *Mach. Learn. Sci. Tech.* 5.3 (2024), p. 035031. DOI: 10.1088/2632-2153/ad66ad. arXiv: 2403.05618 [hep-ph].
- [109] Ting-Hsiang Hsu et al. “EveNet: A Foundation Model for Particle Collision Data Analysis”. In: (Jan. 2026). arXiv: 2601.17126 [hep-ex].
- [110] Farouk Mokhtar et al. “Machine-learned particle flow as a foundation model for collider physics”. In: (June 2026). arXiv: 2606.14373 [hep-ex].
- [111] Supratim Das Bakshi, T. J. Hobbs, and Brandon Kriesten. “Reusable theory representations for colliders: a demonstrator SMEFT foundation model”. In: (Dec. 2025). arXiv: 2512.15862 [hep-ph].
- [112] Liam Parker et al. “AION-1: Omnimodal Foundation Model for Astronomical Sciences”. In: (Oct. 2025). arXiv: 2510.17960 [astro-ph.IM].
- [113] Rikab Gambhir, Luisa Lucie-Smith, and Jesse Thaler. “Interpreting ”Interpretability” and Explaining ”Explainability” in Machine Learning in Physics”. In: (June 2026). arXiv: 2606.26228 [physics.data-an].
- [114] Rikab Gambhir, Athis Osathanapan, and Jesse Thaler. “Streamlining latent spaces in machine learning using moment pooling”. In: *Phys. Rev. D* 110.7 (2024), p. 074020. DOI: 10.1103/PhysRevD.110.074020. arXiv: 2403.08854 [hep-ph].

- [115] Dikshant Sagar et al. “Adapting Vision-Language Models for Neutrino Event Classification in High-Energy Physics”. In: (Sept. 2025). arXiv: 2509.08461 [cs.LG].
- [116] Roberto Franceschini et al. “Kinematic variables and feature engineering for particle phenomenology”. In: *Rev. Mod. Phys.* 95.4 (2023), p. 045004. DOI: 10.1103/RevModPhys.95.045004. arXiv: 2206.13431 [hep-ph].
- [117] Tousif Islam, Digvijay Wadekar, and Zihan Zhou. “gwBenchmarks: Stress-Testing LLM Agents on High-Precision Gravitational Wave Astronomy”. In: (May 2026). arXiv: 2605.11269 [gr-qc].
- [118] Aaron Grattafiori et al. *The Llama 3 Herd of Models*. 2024. arXiv: 2407.21783 [cs.AI]. URL: <https://arxiv.org/abs/2407.21783>.
- [119] OpenAI et al. “gpt-oss-120b and gpt-oss-20b Model Card”. In: (2025). arXiv: 2508.10925 [cs.CL]. URL: <https://arxiv.org/abs/2508.10925>.
- [120] Anthropic. *Model Context Protocol*. <https://modelcontextprotocol.io>. Accessed 2026. 2024.
- [121] Andrew Ng and contributors. *aisuite: Simple, unified interface to multiple generative AI providers*. <https://github.com/andrewyng/aisuite>. GitHub repository, accessed 23 June 2026. 2024.
- [122] Tony Menzo et al. “Reining in an Agentic Harness for High Energy Physics”. In: (Aug. 2026). arXiv: 2609.00107 [hep-ph].